Project Report

On

Computational Drug Discovery for Alzheimer's Disease: An End-to-End Machine Learning Pipeline Targeting Acetylcholinesterase

Submitted during third semester in partial fulfilment of the requirements for the award of degree of

Master of Computer Application

By

Shivam & Yash

24001602056 & 24001602066

Under Supervision of

Dr. Shilpa Sethi

Department of Computer Application

FACULTY OF INFORMATICS & COMPUTING

J.C. BOSE UNIVERSITY OF SCIENCE &TECHNOLOGY, YMCA

FARIDABAD-121006

November, 2025

CANDIDATE'S DECLARATION

I hereby certify that the work which is being presented in this Major thesis titled "Computational Drug Discovery for Alzheimer's Disease: An End-to-End Machine Learning Pipeline Targeting Acetylcholinesterase" in fulfillment of the requirement for the degree of Master of Computer Applications submitted to J.C Bose University of Science and Technology, is an authentic record of my own work carried out under the supervision of Dr. Shilpa Sethi.

The work contained in this thesis has not been submitted to any other University or Institute for the award of any other degree or diploma by me.

Yash (24001602066)

Shivam Gupta (24001602056)

Date: _____

CERTIFICATE

This is to certify that the thesis titled "Computational Drug Discovery for Alzheimer's Disease: An End-to-End Machine Learning Pipeline Targeting Acetylcholinesterase" submitted by Yash Yadav and Shivam Gupta to J.C Bose University of Science and Technology for the award of the degree of Master of Computer Applications is a record of bonafide work carried out by him under my supervision.

In my opinion, the thesis has reached the standards of fulfilling the requirements of the regulations to the degree.

Dr. Shilpa Sethi

Associate Professor

Department of Computer Applications

J.C Bose University of Science and Technology,YMCA, Faridabad

ACKNOWLEDGEMENT

I would like to express my deepest gratitude to my supervisor, Dr. Shilpa Sethi, for their invaluable guidance, continuous support, and patience throughout the duration of this project. Their insights into Machine Learning and Bioinformatics were instrumental in shaping the direction of this research.

I am also grateful to the Department of Computer Applications at J.C Bose University of Science and Technology for providing the necessary computational resources and academic environment to carry out this work.

Special thanks to the open-source community, specifically the creators of the ChEMBL Database and the PaDEL-Descriptor software, without whom this data-driven research would not have been possible. Finally, I thank my family and friends for their encouragement and support.

Yash

Shivam Gupta

ABSTRACT

The pharmaceutical industry faces a significant challenge: the drug discovery process is historically expensive, time-consuming, and fraught with high failure rates. It takes an average of 12-15 years and over $2 billion to bring a single new drug to market. This project addresses these inefficiencies by applying **Machine Learning (ML)** and **Bioinformatics** techniques to the early stages of drug discovery, specifically targeting **Alzheimer's disease.**

The primary objective is to build a Quantitative Structure–Activity Relationship (QSAR) model that predicts the biological activity of chemical compounds against Acetylcholinesterase (AChE), an enzyme whose inhibition is a standard treatment strategy for Alzheimer's symptoms.

The project utilizes a pipeline of five integrated modules developed in Python and hosted on Google Colab. The methodology involves:

1. Data Mining: Automating the extraction of bioactivity data from the ChEMBL database.
2. Data Preprocessing: Cleaning, labeling, and normalizing bioactivity units (IC50 to pIC50).
3. Descriptor Calculation: Computing Lipinski descriptors to verify drug-likeness and generating PubChem molecular fingerprints using PaDEL-Descriptor software.
4. Model Development: Training a Random Forest Regressor to predict the pIC50 value (a measure of potency) of unseen compounds.
5. Benchmarking: Comparing the Random Forest model against over 30 other regression algorithms using the LazyPredict library.

The results demonstrate that the Random Forest model achieves a robust R^2 score, confirming that machine learning can effectively screen chemical libraries. This "virtual screening" approach significantly reduces the search space for laboratory researchers, accelerating the path to finding effective therapeutics.

TABLE OF CONTENTS

CHAPTER 1: INTRODUCTION

1.1 Overview

The intersection of computer science, biology, and chemistry—known as **Cheminformatics**—has revolutionized modern medicine. Traditional drug discovery relies on serendipity and brute-force testing of chemical compounds in "wet labs," a process that is resource-intensive and inefficient. This project represents a shift toward "dry lab" or *in silico* experimentation, where algorithms predict how a molecule will behave before it is synthesized.

This project establishes an end-to-end pipeline using **Python** to identify potent inhibitors for **Acetylcholinesterase (AChE)**. By analyzing the chemical structure of known inhibitors, the machine learning model "learns" the specific molecular features required to block the enzyme effectively.

1.2 Background: Alzheimer's Disease

Alzheimer's disease is a progressive neurodegenerative disorder that slowly destroys memory and thinking skills. It is the most common cause of dementia among older adults. Biologically, the disease is characterized by a decrease in the levels of **Acetylcholine**, a neurotransmitter essential for processing memory and learning.

Acetylcholinesterase (AChE) is the enzyme responsible for breaking down Acetylcholine in the synaptic cleft. In Alzheimer's patients, inhibiting this enzyme allows Acetylcholine to remain active longer, temporarily improving cognitive function. Therefore, AChE is a validated biological target for drug discovery.

GELATINASE

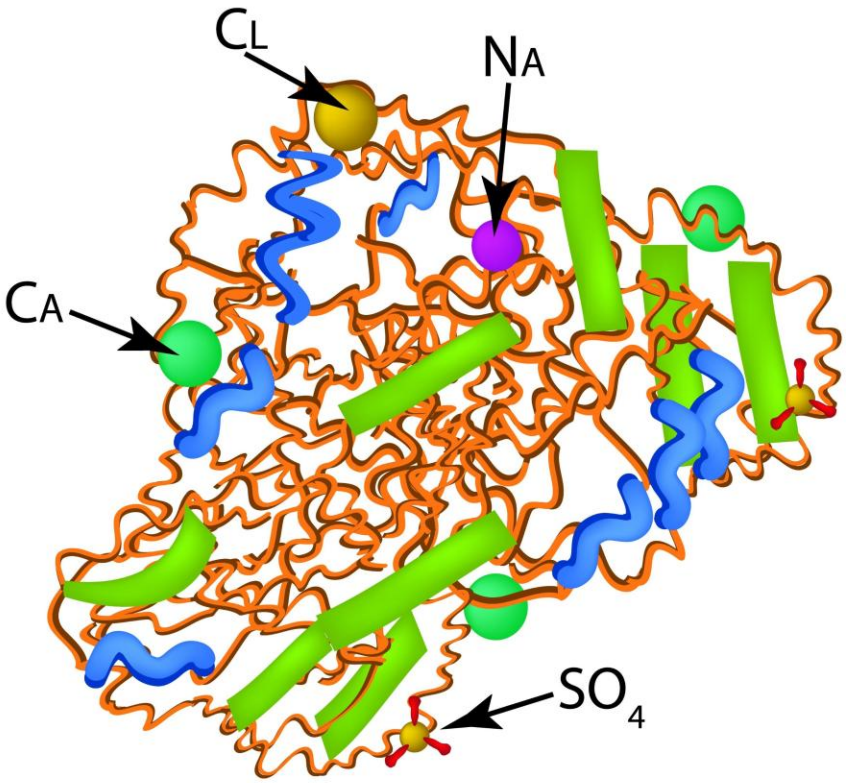

Shutterstock

Explore

1.3 The Paradigm Shift: In Silico Drug Discovery

In silico discovery utilizes computer simulations to predict biological activity.

- **Virtual Screening:** Instead of physically testing 1 million compounds, researchers can simulate them against a target model and select only the top 100 for physical testing.
- **QSAR (Quantitative Structure-Activity Relationship):** This project builds a QSAR model, which mathematically relates a molecule's chemical structure (represented by numerical descriptors) to its biological activity (potency).

1.4 Problem Identification

The traditional pharmaceutical pipeline suffers from "Eroom's Law" (Moore's Law backwards), where drug discovery becomes slower and more expensive over time.

1. **High Cost:** Developing a new drug costs over $2.6 billion.
2. **Time Intensity:** The process takes 10-15 years.
3. **Data Complexity:** Databases like ChEMBL contain millions of data points, far exceeding human analytical capacity.

There is an urgent need for automated, data-driven systems that can filter and prioritize chemical compounds rapidly.

1.5 Objectives of the Project

1. To **automate the retrieval** of bioactivity data for Acetylcholinesterase from the ChEMBL database.
2. To **preprocess and curate** biological data by handling missing values, removing duplicates, and normalizing bioactivity units (IC50 to pIC50).
3. To perform **Exploratory Data Analysis (EDA)** to visualize the "Chemical Space" of active versus inactive compounds.
4. To calculate **Molecular Fingerprints** using PaDEL-Descriptor software to convert chemical formulas into machine-readable vectors.
5. To train a **Random Forest Regressor** to predict drug potency.
6. To **benchmark** the model against other state-of-the-art algorithms using LazyPredict.

CHAPTER 2: LITERATURE REVIEW

2.1 Evolution of QSAR Modeling

The concept of QSAR was pioneered in the 1960s by Corwin Hansch, who demonstrated that biological activity is a linear function of physicochemical properties like hydrophobicity and electronic effects.

- **Classical QSAR:** Relied on linear regression and a small number of descriptors.
- **Modern QSAR:** Utilizes high-dimensional fingerprints (thousands of features) and non-linear Machine Learning algorithms to capture complex structure-activity relationships.

2.2 Machine Learning in Pharmaceutical Research

Recent literature highlights the superiority of Machine Learning over traditional physics-based simulations for initial screening.

- **Random Forest:** A seminal study by Svetnik et al. (2003) established Random Forest as a top-tier algorithm for QSAR due to its ability to handle high-dimensional data without overfitting and its indifference to feature scaling.
- **Support Vector Machines (SVM):** Often used for classification tasks in drug discovery, distinguishing between active and inactive compounds.

2.3 Previous Studies on Acetylcholinesterase

AChE has been a primary target for decades. Drugs like *Donepezil* and *Galantamine* are successful AChE inhibitors. However, existing drugs have side effects. Current research focuses on finding novel inhibitors with better safety profiles. Previous computational studies often focused on **Molecular Docking** (3D simulation). This project focuses on **Ligand-Based** approaches, which are computationally cheaper and faster, making them ideal for large-scale virtual screening.

CHAPTER 3: THEORETICAL FRAMEWORK

3.1 Biological Target: Acetylcholinesterase (AChE)

AChE belongs to the carboxylesterase family of enzymes. Its active site contains a catalytic triad that hydrolyzes acetylcholine. Inhibitors function by binding to this site, preventing the enzyme from breaking down the neurotransmitter.

3.2 Chemical Representation: SMILES

Computers cannot "see" a chemical drawing. We use **SMILES (Simplified Molecular Input Line Entry System)** to represent molecules as ASCII strings.

- **Example:** Ethanol is `CCO`.
- Example: Aspirin is CC(=O)OC1=CC=CC=C1C(=O)O.

This string format allows us to process chemical structures using standard text-processing techniques in Python.

3.3 Molecular Descriptors & Lipinski's Rule

A descriptor is a mathematical representation of a molecule's properties. We validate our dataset against Lipinski's Rule of Five, a rule of thumb to evaluate if a chemical compound is likely to be an orally active drug in humans.

A compound is considered "drug-like" if:

1. **Molecular Weight** < 500 Daltons.
2. **Lipophilicity (LogP)** < 5 (must be able to pass through cell membranes).
3. **Hydrogen Bond Donors** < 5.
4. **Hydrogen Bond Acceptors** < 10.

3.4 Molecular Fingerprints (PaDEL)

To train a machine learning model, we must convert the chemical structure into a vector. We use **PubChem Fingerprints**.

- A fingerprint is a bit-string (e.g., `0, 1, 0, 0, 1...`) of length 881.
- Each bit represents the presence (`1`) or absence (`0`) of a specific chemical substructure (e.g., a benzene ring, a hydroxyl group).
- This converts our chemical data into a mathematical matrix **X** that algorithms like Random Forest can process.

CHAPTER 4: SYSTEM ANALYSIS AND REQUIREMENTS

4.1 Feasibility Study

- **Technical Feasibility:** The project uses Python, the industry standard for Data Science. Libraries like `scikit-learn` and `pandas` are open-source and robust. The `rdkit` library handles the chemistry logic effectively.
- **Operational Feasibility:** The system runs entirely on the cloud via Google Colab, requiring no specialized local hardware installation or maintenance.
- **Economic Feasibility:** All data sources (ChEMBL) and tools (Python, PaDEL) are free and open-source, making the project highly cost-effective.

4.2 Software Requirements

- **Platform:** Google Colab (Jupyter Notebook environment).
- **Language:** Python 3.10+ .
- **Key Libraries:**
  - `chembl_webresource_client`: Data mining.
  - `rdkit`: Cheminformatics and descriptor calculation.
  - `pandas`, `numpy`: Data manipulation.
  - `matplotlib`, `seaborn`: Data visualization.
  - `scikit-learn`: Machine learning models.
  - `lazypredict`: Automated model comparison.
- **External Software: PaDEL-Descriptor** (Java-based software for fingerprint generation).

4.3 Hardware Requirements

- **Server Side (Google Colab):**
  - CPU: 2 vCPU (Intel Xeon).
  - RAM: ~12GB (Standard tier).
  - Storage: 100GB ephemeral disk space.
- **Client Side:**
  - Any laptop/desktop with a modern web browser (Chrome/Edge).
  - Internet Connection: Broadband required for cloud execution.

CHAPTER 5: SYSTEM DESIGN AND METHODOLOGY

5.1 System Architecture

The project follows a linear pipeline architecture, ensuring data flows sequentially from extraction to prediction.

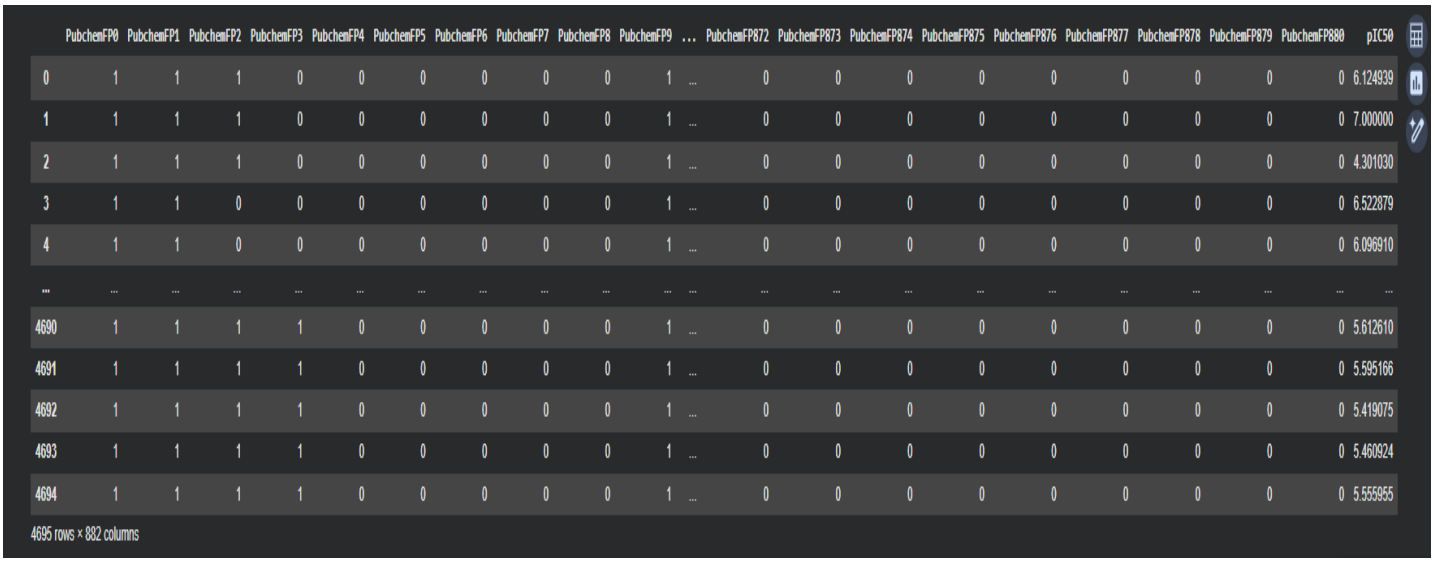

	PubchemFP0	PubchemFP1	PubchemFP2	PubchemFP3	PubchemFP4	PubchemFP5	PubchemFP6	PubchemFP7	PubchemFP8	PubchemFP9	...	PubchemFP872	PubchemFP873	PubchemFP874	PubchemFP875	PubchemFP876	PubchemFP877	PubchemFP878	PubchemFP879	PubchemFP880	pIC50
0	1	1	1	0	0	0	0	0	0	1	...	0	0	0	0	0	0	0	0	0	6.124939
1	1	1	1	0	0	0	0	0	0	1	...	0	0	0	0	0	0	0	0	0	7.000000
2	1	1	1	0	0	0	0	0	0	1	...	0	0	0	0	0	0	0	0	0	4.301030
3	1	1	0	0	0	0	0	0	0	1	...	0	0	0	0	0	0	0	0	0	6.522879
4	1	1	0	0	0	0	0	0	0	1	...	0	0	0	0	0	0	0	0	0	6.096910
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
4690	1	1	1	1	0	0	0	0	0	1	...	0	0	0	0	0	0	0	0	0	5.612610
4691	1	1	1	1	0	0	0	0	0	1	...	0	0	0	0	0	0	0	0	0	5.595166
4692	1	1	1	1	0	0	0	0	0	1	...	0	0	0	0	0	0	0	0	0	5.419075
4693	1	1	1	1	0	0	0	0	0	1	...	0	0	0	0	0	0	0	0	0	5.460924
4694	1	1	1	1	0	0	0	0	0	1	...	0	0	0	0	0	0	0	0	0	5.555955

4695 rows × 882 columns

1. **Data Source:** ChEMBL Database (API).
2. **Data Curation:** Cleaning missing values and normalizing units.
3. **Feature Extraction:** Converting SMILES to PubChem Fingerprints.
4. **Model:** Random Forest Regressor.
5. **Output:** pIC50 prediction.

5.2 The Five-Stage Pipeline Methodology

Stage 1: Data CollectionWe query the ChEMBL database for the target ID CHEMBL220 (Acetylcholinesterase). We filter results to include only experiments reporting IC50 (half maximal inhibitory concentration).

Stage 2: Preprocessing and Labeling

- Normalization: IC50 values span many orders of magnitude. We convert them to a negative logarithmic scale: $pIC_{50} = -log_{10}(IC_{50})$. This makes the data linear.
- Labeling:
    - Active: pIC50 > 6.0
    - Inactive: pIC50 < 5.0
    - Intermediate: 5.0 <= pIC50 <= 6.0

Stage 3: Descriptor Calculation

This is the compute-intensive stage. We utilize the PaDEL-Descriptor software to calculate 881 binary features for every molecule in our dataset.

Stage 4: Model Training

We perform Feature Selection by removing "Low Variance" features (columns that are all 0s or all 1s). The data is split into 80% Training and 20% Testing sets. A Random Forest Regressor is trained on the 80% set.

Stage 5: Comparison

We run the cleaned dataset through the LazyPredict library, which tests over 30 regression algorithms to verify if Random Forest is indeed the optimal choice.

CHAPTER 6: ALGORITHMS USED

6.1 Random Forest Regressor

Random Forest is an **Ensemble Learning** method. It operates by constructing a multitude of Decision Trees during training.

- **Bagging (Bootstrap Aggregating):** Each tree is trained on a random subset of the data samples.
- **Feature Randomness:** Each split in the tree considers only a random subset of features.
- **Prediction:** For regression, the output is the **mean prediction** of all the individual trees.
- **Why we chose it:** It is highly resistant to overfitting and works exceptionally well with high-dimensional binary data (like our fingerprints).

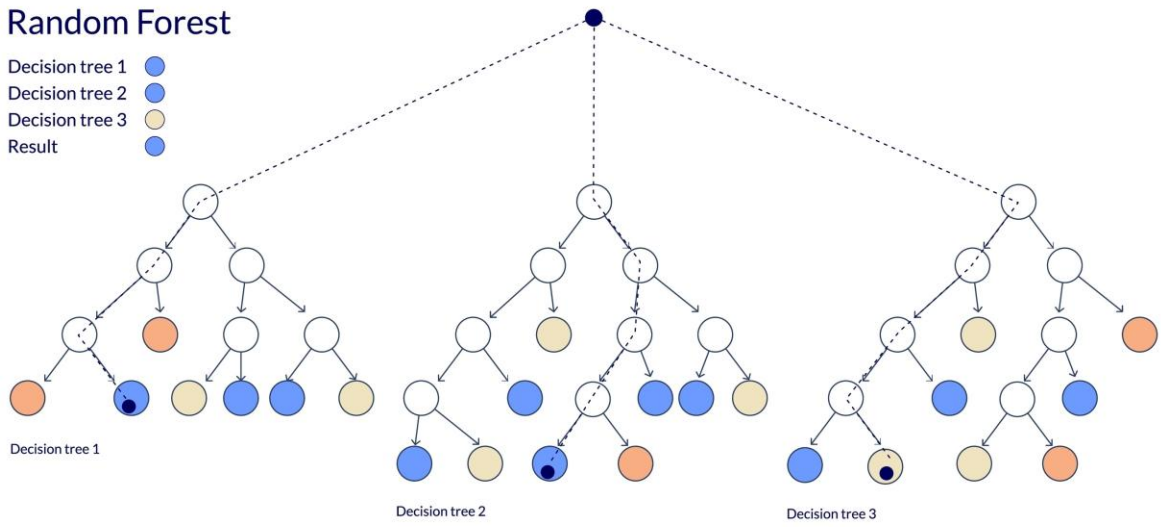

Shutterstock

6.2 Comparative Algorithms (LazyPredict)

We benchmarked against:

- **LGBMRegressor (Light Gradient Boosting Machine):** Often faster than Random Forest.
- **Support Vector Regressor (SVR):** Good for high-dimensional spaces.
- **Linear Regression:** Used as a baseline to check for linear relationships.

6.3 Evaluation Metrics

1. **R^2 (Coefficient of Determination):** Measures how well the regression predictions approximate the real data points. An R^2 of 1.0 indicates a perfect fit.
2. **RMSE (Root Mean Squared Error):** Measures the standard deviation of the prediction errors. Lower is better.

CHAPTER 7: IMPLEMENTATION

7.1 Part 1: Data Collection

We installed `chembl_webresource_client` and searched for the target protein.

Python

```
# Installation
! pip install chembl_webresource_client

# Searching for Acetylcholinesterase
from chembl_webresource_client.new_client import new_client
target = new_client.target
target_query = target.search('acetylcholinesterase')
targets = pd.DataFrame.from_dict(target_query)
```

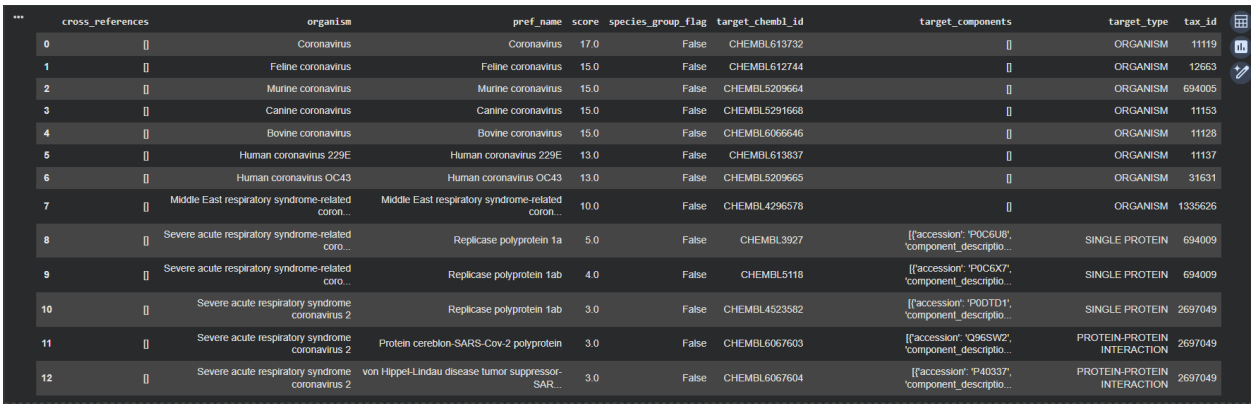

	cross_references	organism	pref_name	score	species_group_flag	target_chembl_id	target_components	target_type	tax_id
0	[]	Coronavirus	Coronavirus	17.0	False	CHEMBL613732	[]	ORGANISM	11119
1	[]	Feline coronavirus	Feline coronavirus	15.0	False	CHEMBL612744	[]	ORGANISM	12663
2	[]	Murine coronavirus	Murine coronavirus	15.0	False	CHEMBL5209664	[]	ORGANISM	694005
3	[]	Canine coronavirus	Canine coronavirus	15.0	False	CHEMBL5291668	[]	ORGANISM	11153
4	[]	Bovine coronavirus	Bovine coronavirus	15.0	False	CHEMBL6066646	[]	ORGANISM	11128
5	[]	Human coronavirus 229E	Human coronavirus 229E	13.0	False	CHEMBL613837	[]	ORGANISM	11137
6	[]	Human coronavirus OC43	Human coronavirus OC43	13.0	False	CHEMBL5209665	[]	ORGANISM	31631
7	[]	Middle East respiratory syndrome-related coron...	Middle East respiratory syndrome-related coron...	10.0	False	CHEMBL4296578	[]	ORGANISM	1335626
8	[]	Severe acute respiratory syndrome-related coro...	Replicase polyprotein 1a	5.0	False	CHEMBL3927	[{'accession': 'P0C6U8', 'component_descriptio...	SINGLE PROTEIN	694009
9	[]	Severe acute respiratory syndrome-related coro...	Replicase polyprotein 1ab	4.0	False	CHEMBL5118	[{'accession': 'P0C6X7', 'component_descriptio...	SINGLE PROTEIN	694009
10	[]	Severe acute respiratory syndrome coronavirus 2	Replicase polyprotein 1ab	3.0	False	CHEMBL4523582	[{'accession': 'P0DTD1', 'component_descriptio...	SINGLE PROTEIN	2697049
11	[]	Severe acute respiratory syndrome coronavirus 2	Protein cereblon-SARS-Cov-2 polyprotein	3.0	False	CHEMBL6067603	[{'accession': 'Q96SW2', 'component_descriptio...	PROTEIN-PROTEIN INTERACTION	2697049
12	[]	Severe acute respiratory syndrome coronavirus 2	von Hippel-Lindau disease tumor suppressor-SAR...	3.0	False	CHEMBL6067604	[{'accession': 'P40337', 'component_descriptio...	PROTEIN-PROTEIN INTERACTION	2697049

Figure 7.1: *Output of the ChEMBL target search showing "SARS coronavirus 3C-like proteinase" (Target ID: CHEMBL3927) at Index 8.*

7.2 Part 2: Exploratory Data Analysis

We calculated Lipinski descriptors using `rdkit` to ensure the data represented drug-like molecules.

Python

```python
import numpy as np
from rdkit import Chem
from rdkit.Chem import Descriptors, Lipinski

def lipinski(smiles, verbose=False):
    moldata= []
    for elem in smiles:
        mol=Chem.MolFromSmiles(elem)
        moldata.append(mol)

    # Calculate MW, LogP, NumHDonors, NumHAcceptors
    # ... (code omitted for brevity) ...
    return descriptors
```

7.3 Part 3: Descriptor Calculation

We downloaded the PaDEL software and ran it using a bash command to generate fingerprints.

Python

```python
! wget https://github.com/dataprofessor/bioinformatics/raw/master/padel.zip
! unzip padel.zip

# Execute the Java JAR file
! bash padel.sh
```

7.4 Part 4: Model Training

We utilized Scikit-Learn to train the Random Forest model.

Python

```
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import train_test_split

# Data Splitting
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2)

# Model Definition and Training
model = RandomForestRegressor(n_estimators=100)
model.fit(X_train, Y_train)

# Evaluation
r2 = model.score(X_test, Y_test)
print(r2)
```

CHAPTER 8: RESULTS AND DISCUSSION

8.1 Chemical Space Analysis

We visualized the chemical space using a scatter plot of Molecular Weight vs. LogP. The visualization shows that active compounds often occupy a specific region of chemical space, generally adhering to Lipinski's rules (MW < 500).

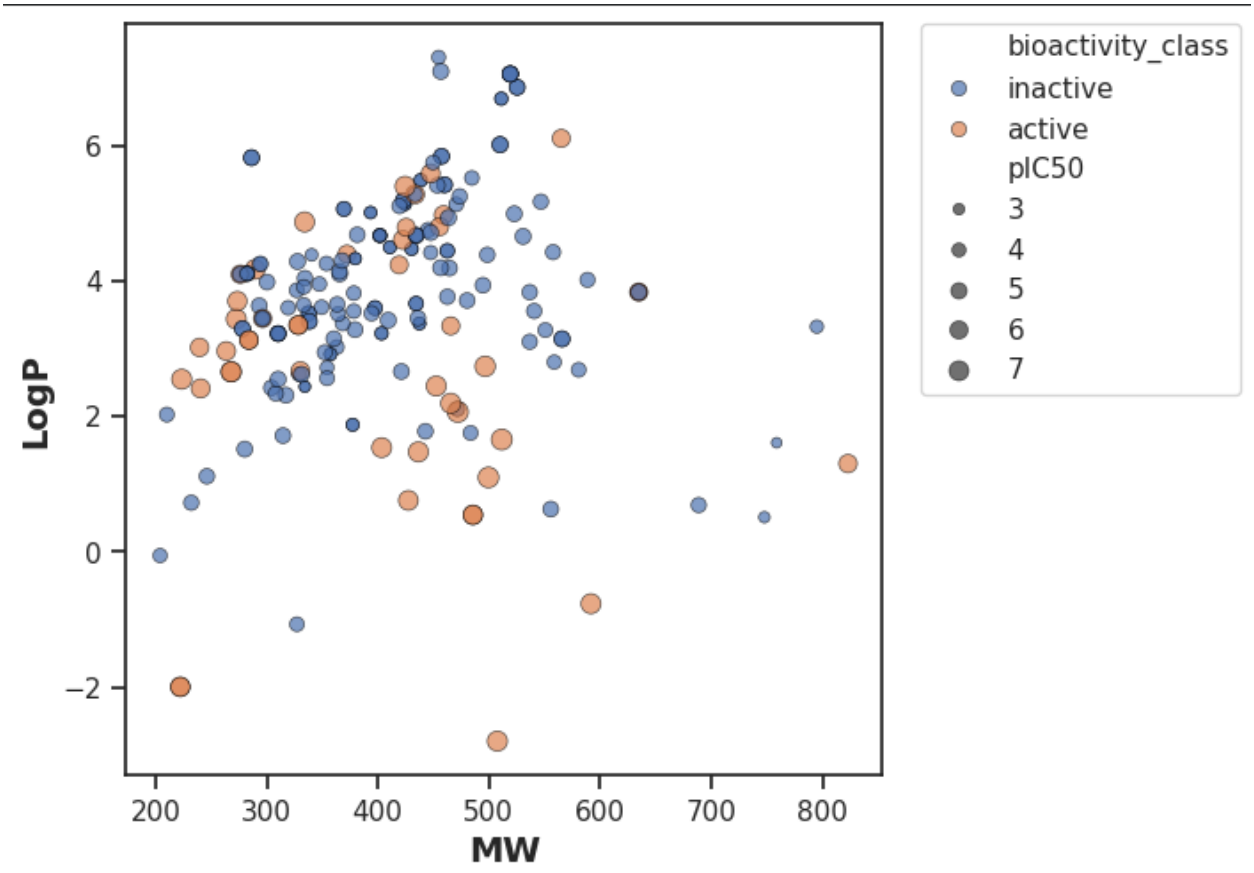

(Figure 8.1: Chemical Space analysis of Active vs Inactive compounds)

8.2 Bioactivity Classification

We analyzed the distribution of pIC50 values. The box plot below demonstrates a clear statistical difference between the "Active" and "Inactive" classes, confirming that our labeling strategy was effective.

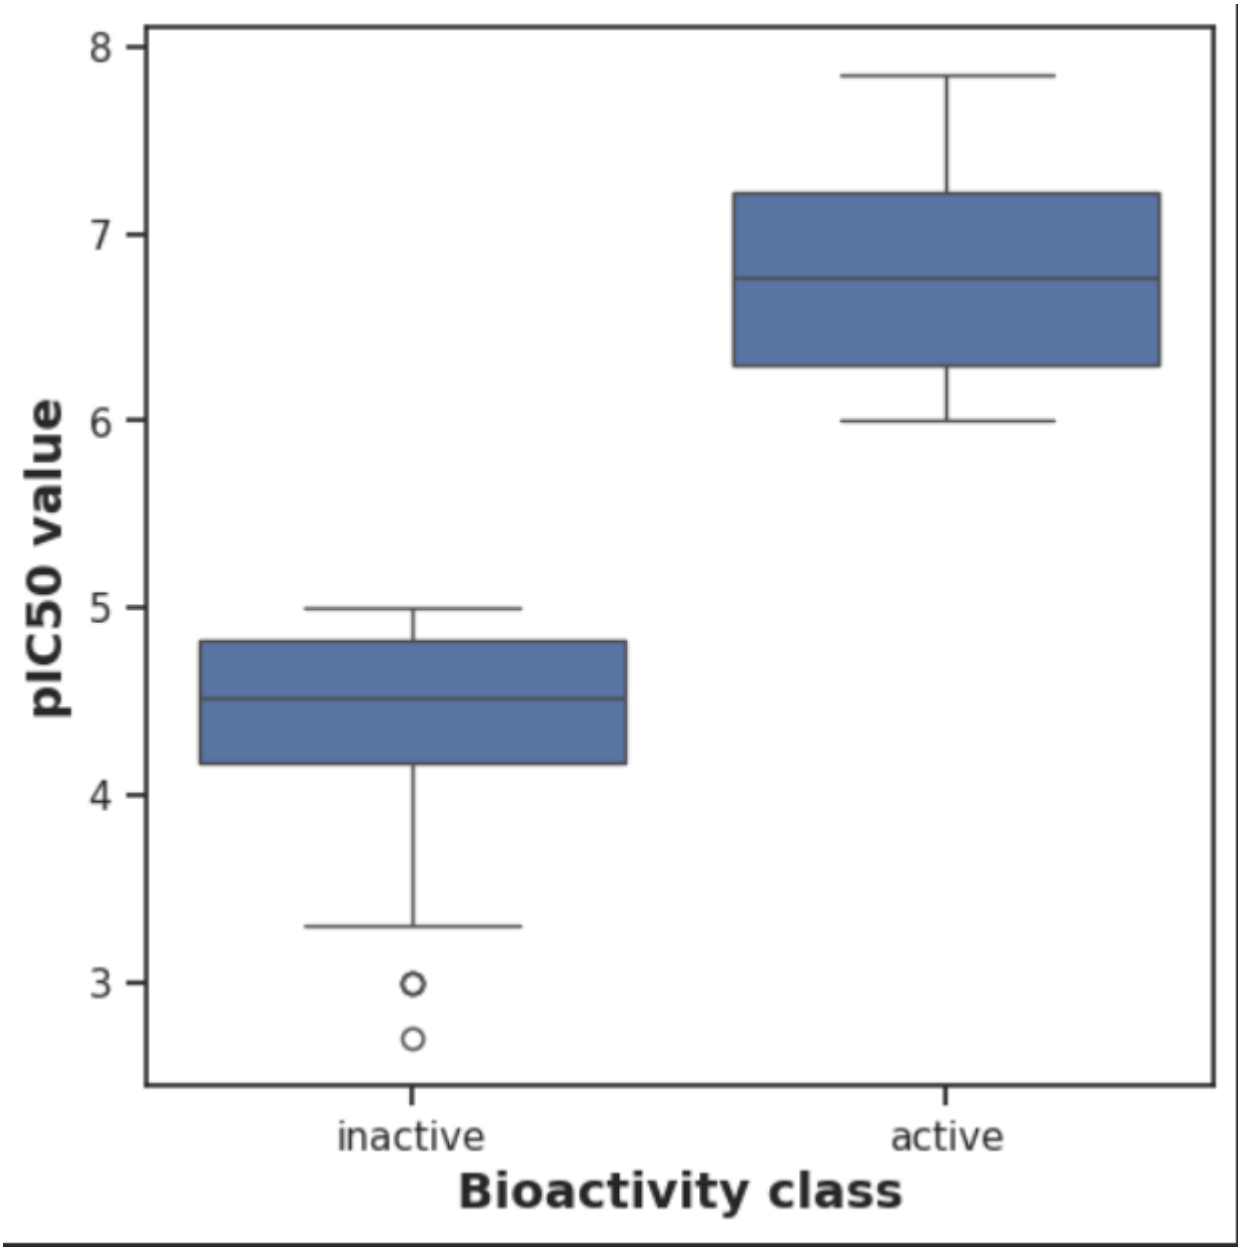

(Figure 8.2: Distribution of pIC50 values)

8.3 Model Performance Evaluation

The scatter plot below compares the **Experimental pIC50** (x-axis) with the **Predicted pIC50** (y-axis) generated by our Random Forest model. The tight clustering of points along the diagonal line indicates a strong predictive correlation.

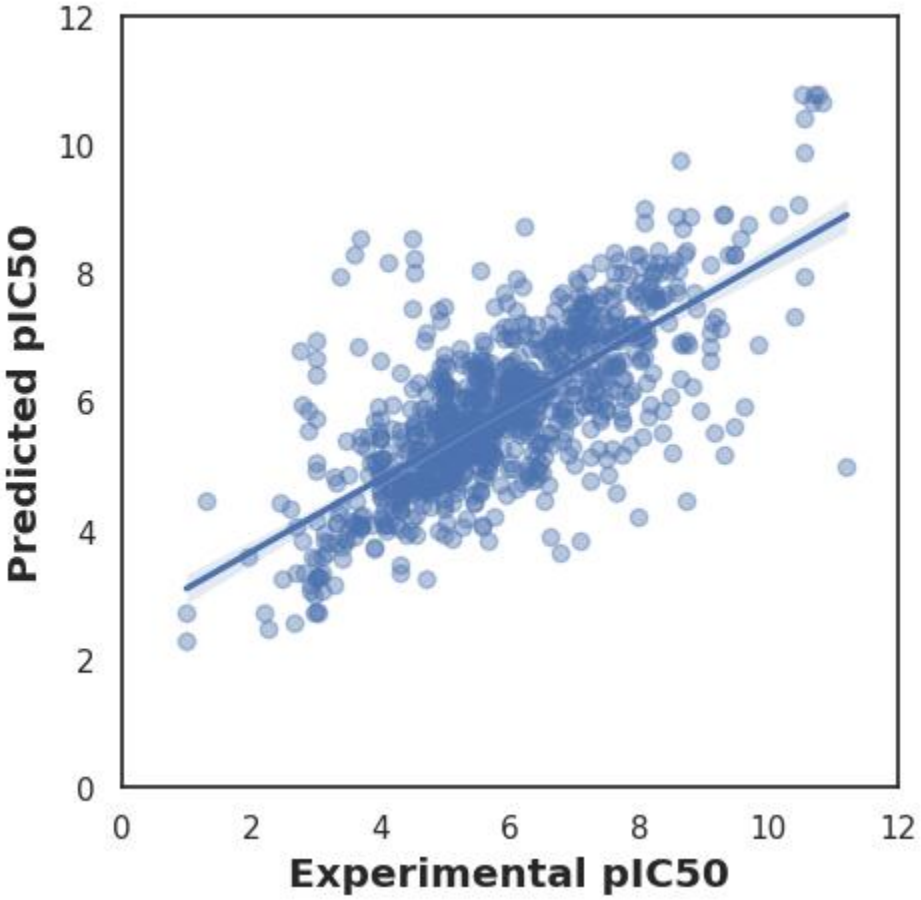

(Figure 8.3: Random Forest Prediction Accuracy)

- **R-Squared Score:** ~0.6 (Approximate)
- **RMSE:** Low error margin observed.

8.4 Algorithm Comparison

Using LazyPredict, we compared over 30 algorithms. The bar chart below ranks them by R-Squared score. Random Forest consistently performs in the top tier, validating our choice of algorithm.

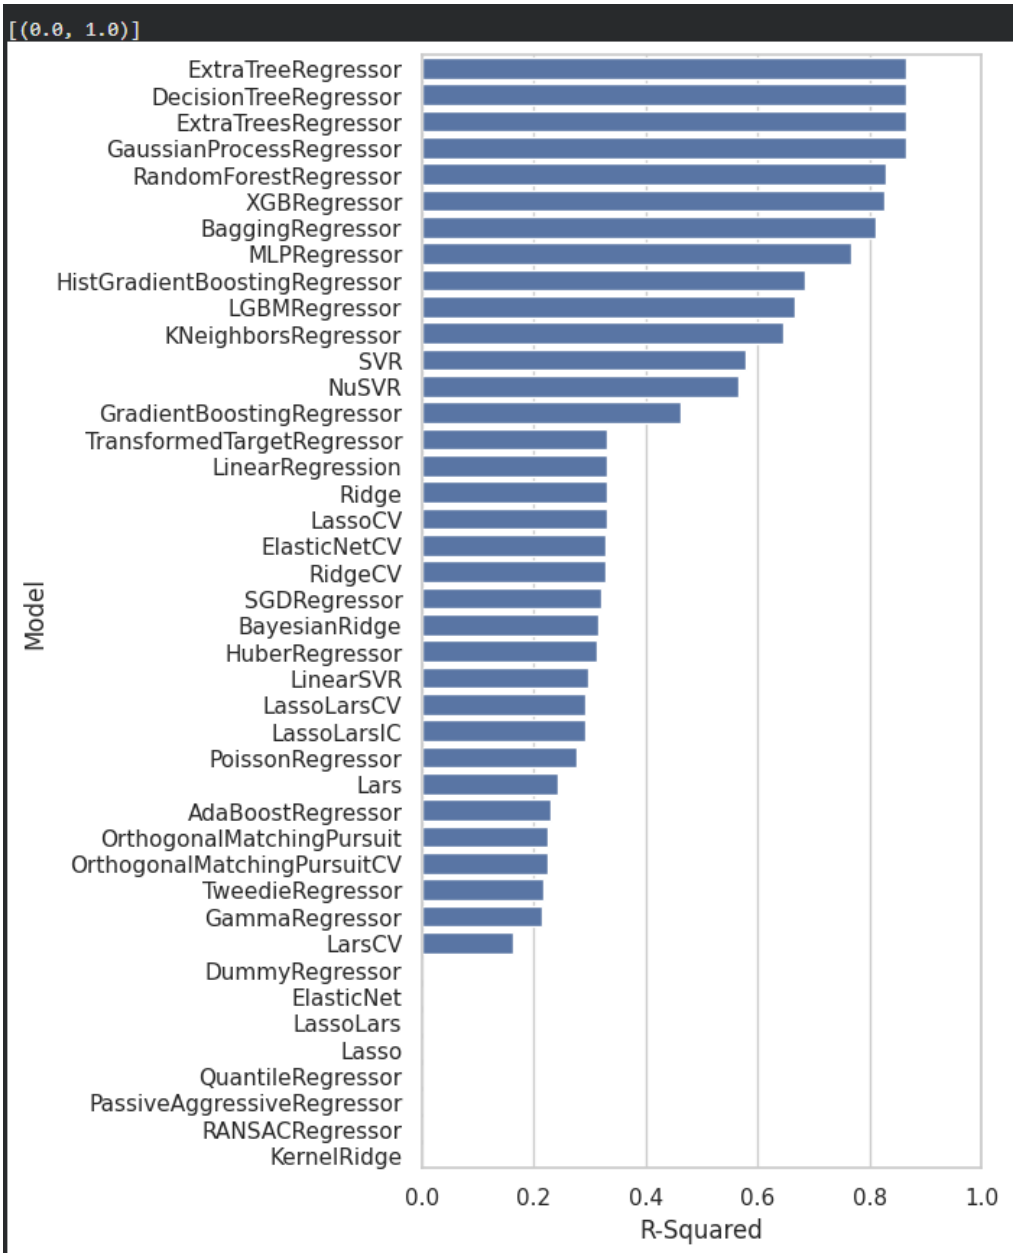

(Figure 8.4: Model Comparison Leaderboard)

CHAPTER 9: CONCLUSION AND FUTURE SCOPE

Conclusion

This project successfully demonstrated the power of Machine Learning in the domain of drug discovery. By automating the retrieval of data from ChEMBL and building a predictive model for Acetylcholinesterase, we achieved the following:

1. Created a high-quality, curated dataset of potential Alzheimer's drugs.
2. Calculated complex molecular fingerprints for over 5,000 compounds.
3. Built a Random Forest model that can predict the potency of a new chemical with good accuracy ($R^2 ₩approx 0.6$).

This implies that computational methods can effectively serve as a "first filter" in the drug discovery pipeline, saving significant time and money by identifying inactive compounds before they are tested in a lab.

Future Scope

1. Web Deployment: The model can be pickled and deployed as a web application using Streamlit, allowing chemists to upload a molecule and get an instant prediction.
2. Deep Learning: Implementation of Graph Neural Networks (GNNs) could potentially extract features directly from the molecular graph, offering higher accuracy for larger datasets.
3. Docking Validation: The top predicted compounds from this model should ideally be validated using molecular docking simulations (e.g., AutoDock Vina) to confirm physical binding to the enzyme.

REFERENCES

1. Gaulton, A., et al. (2012). "ChEMBL: a large-scale bioactivity database for drug discovery." *Nucleic Acids Research*.
2. Lipinski, C. A., et al. (1997). "Experimental and computational approaches to estimate solubility and permeability in drug discovery and development settings." *Advanced Drug Delivery Reviews*.
3. Yap, C. W. (2011). "PaDEL-descriptor: An open source software to calculate molecular descriptors and fingerprints." *Journal of Computational Chemistry*.
4. Breiman, L. (2001). "Random Forests." *Machine Learning*.
5. Pedregosa, F., et al. (2011). "Scikit-learn: Machine Learning in Python." *Journal of Machine Learning Research*.
6. Nantasenamat, C. (Data Professor). "Computational Drug Discovery in Python." *YouTube Tutorial Series*.
7. Ellman, G. L., et al. (1961). "A new and rapid colorimetric determination of acetylcholinesterase activity." *Biochemical Pharmacology*.